

PA4_EE21S056

October 30, 2022

Akash Sharma (EE21S056)

EE5179: Deep Learning for Imaging

1 Comparing PCA and Autoencoders

1.1 Importing packages

```
[1]: import torch
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor, Lambda
from torch import nn
from torchvision.utils import make_grid
import matplotlib.pyplot as plt
import numpy as np
```

1.2 Loading MNIST dataset

```
[2]: training_data = datasets.MNIST(root="data", train=True, download =
    ↳True, transform = ToTensor())
test_data = datasets.MNIST(root="data", train = False, download =
    ↳True, transform = ToTensor())

batch_size = 64

train_loader = DataLoader(dataset=training_data, batch_size=batch_size,
    ↳shuffle=True)
test_loader = DataLoader(dataset=test_data, batch_size=batch_size,
    ↳shuffle=False)
```

1.3 MNIST Image Reconstruction using PCA

Reference used: <https://analyticsindiamag.com/guide-to-image-reconstruction-using-principal-component-analysis/>

```
[3]: # Extracting just images from the dataset for PCA
```

```
images = training_data.data  
print(images.shape)
```

```
torch.Size([60000, 28, 28])
```

1.3.1 Finding indices representing each class

```
[4]: # 114-0, 102-1, 117-2, 111-3, 115-4, 145-5, 106-6, 103-7, 125-8, 110-9
```

```
idx_PCA = [114, 102, 117, 111, 115, 145, 106, 103, 125, 110]
```

```
i=1 #just an iterator variable
```

```
for idx in idx_PCA:  
    plt.subplot(2,5,i)  
    plt.imshow(images[idx].reshape(28,28), cmap='gray')  
    plt.axis('off')  
    i+=1  
plt.tight_layout()  
plt.show()
```



1.3.2 Flattening PCA input images

```
[5]: data = np.zeros((60000, 784))  
k=0  
for img in images:  
    data[k] = img.squeeze().flatten()
```

```

data[k] = data[k]/255
k+=1
print(data.shape)

```

(60000, 784)

1.3.3 Performing PCA and checking reconstruction error

```

[6]: from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error
pca = PCA(n_components=30)
reconst_data = pca.inverse_transform(pca.fit_transform(data))

i = 1

for idx in idx_PCA:
    plt.subplot(2,5,i)
    plt.imshow(reconst_data[idx].reshape(28,28), cmap='gray')
    plt.axis('off')
    i+=1

plt.tight_layout()
plt.show()

print('Reconstruction error for PCA is ', mean_squared_error(data,
↪reconst_data))

```



Reconstruction error for PCA is 0.018121850503857066

1.4 Using an Autoencoder on MNIST dataset

1.4.1 Defining stacked autoencoder class

```
[7]: class stack_AE(nn.Module):
    def __init__(self):
        super(stack_AE, self).__init__()
        self.flatten = nn.Flatten()
        self.encoder = nn.Sequential(
            nn.Linear(784,512),
            nn.ReLU(),
            nn.Linear(512,256),
            nn.ReLU(),
            nn.Linear(256,128),
            nn.ReLU(),
            nn.Linear(128,30),
            nn.ReLU())
        self.decoder =nn.Sequential(
            nn.Linear(30,128),
            nn.ReLU(),
            nn.Linear(128,256),
            nn.ReLU(),
            nn.Linear(256,784),
            nn.ReLU())

    def forward(self, x):
        x = self.flatten(x)
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded, encoded
```

1.4.2 Autoencoder Training

```
[8]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

learning_rate = 0.001
epochs = 10

criterion = nn.MSELoss() # Reconstruction error is defined as the MSE Loss

model_stackAE = stack_AE().to(device)
optimizer = torch.optim.Adam(model_stackAE.parameters(), learning_rate)

train_batch_loss = []
train_epoch_loss = []
val_loss = []
```

```

[9]: for epoch in range(epochs):
    for batch_idx, (data, target) in enumerate(train_loader):
        data = data.to(device)
        target = target.to(device)

        model_stackAE.train()
        decoded, encoded = model_stackAE(data)
        loss = criterion(decoded, torch.flatten(data,1))
        train_batch_loss.append(loss.item())

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    train_epoch_loss.append(train_batch_loss[-1])

    model_stackAE.eval()
    with torch.no_grad():
        for (data, target) in test_loader:
            data = data.to(device)
            target = target.to(device)
            decoded, encoded = model_stackAE(data)
            val_iter_loss = criterion(decoded, torch.flatten(data,1))

    val_loss.append(val_iter_loss.item())
    print('Epoch Number: ', epoch+1)
    print('Training loss: ', train_epoch_loss[-1])
    print('Validation loss: ', val_loss[-1])
    print('-----')

print('Training done!')

```

```

Epoch Number:  1
Training loss:  0.02948594279587269
Validation loss:  0.025837989524006844
-----

```

```

Epoch Number:  2
Training loss:  0.02277982421219349
Validation loss:  0.020370956510305405
-----

```

```

Epoch Number:  3
Training loss:  0.01931915245950222
Validation loss:  0.018611568957567215
-----

```

```

Epoch Number:  4
Training loss:  0.02137576974928379
Validation loss:  0.017161209136247635
-----

```

```

Epoch Number: 5
Training loss: 0.018504636362195015
Validation loss: 0.01636418327689171
-----
Epoch Number: 6
Training loss: 0.020677324384450912
Validation loss: 0.01690388098359108
-----
Epoch Number: 7
Training loss: 0.016058238223195076
Validation loss: 0.016165537759661674
-----
Epoch Number: 8
Training loss: 0.015046187676489353
Validation loss: 0.015822740271687508
-----
Epoch Number: 9
Training loss: 0.016513552516698837
Validation loss: 0.01591390371322632
-----
Epoch Number: 10
Training loss: 0.016025178134441376
Validation loss: 0.015233645215630531
-----
Training done!

```

1.4.3 Visualizing training and validation loss for autoencoder

```

[10]: fig, axs = plt.subplots(2, constrained_layout= True, figsize=(8,8))
fig.suptitle('Plots for Stacked Autoencoder')
axs[0].set(ylabel='Training Loss', xlabel='Epochs')
axs[0].plot(range(1,epochs+1),train_epoch_loss, 'tab:red')
axs[1].set(ylabel='Validation Loss', xlabel='Epochs')
axs[1].plot(range(1,epochs+1),val_loss, 'tab:green')


```

```

[10]: [<matplotlib.lines.Line2D at 0x7fe36254f710>]

```

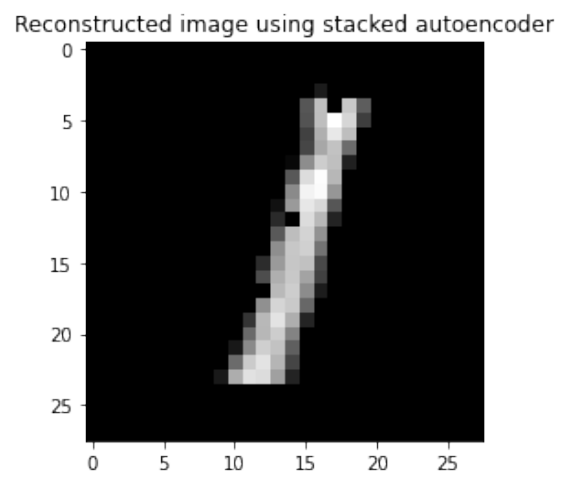

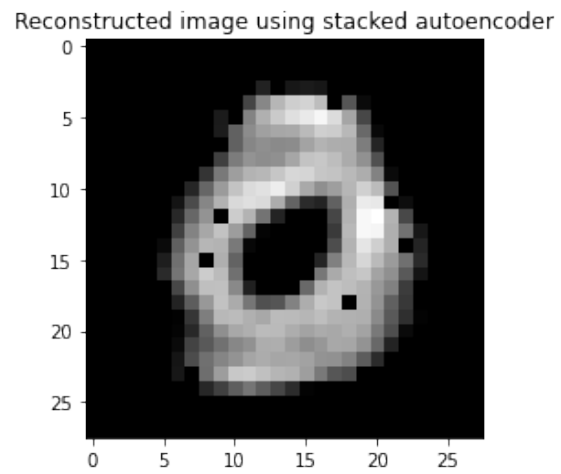

Plots for Stacked Autoencoder















```
[11]: idx_test=[3,2,1,18,4,8,11,0,84,7]





for idx in idx_test:
    image = test_loader.dataset.data[idx, :, :].clone()
    with torch.no_grad():
        image = image.view(-1,28,28).to(device=device).float()
        decoded, encoded = model_stackAE.forward(image)
        decoded = decoded.detach().cpu().numpy()
        image = image.reshape(28,28).detach().cpu().numpy()
        plt.subplot(1,2,1)
        plt.imshow(image , cmap='gray')
        plt.title('input image')
        plt.subplots_adjust(right=1.5)
```

```
plt.subplot(1,2,2)
plt.imshow(decoded.reshape(28,28),cmap = 'gray')
plt.title("Reconstructed image using stacked autoencoder")
plt.show()
```









```
[12]: model_stackAE.eval()
test_loss = 0
with torch.no_grad():
    for (data, target) in test_loader:
        data = data.to(device)
        target = target.to(device)
        decoded, encoded = model_stackAE(data)
        val_iter_loss = criterion(decoded, torch.flatten(data,1))
        test_loss += val_iter_loss/len(test_loader)
```

```
print('Reconstruction accuracy with the given autoencoder configuration is ',  
      ↪test_loss.item())
```

Reconstruction accuracy with the given autoencoder configuration is
0.01669379137456417

2 Experimenting with hidden units of varying sizes

2.1 Defining standard autoencoder class

The following model takes the size of hidden layer as input as required by the experiment.

```
[13]: class standard_AE(nn.Module):  
    def __init__(self,hidden_layer):  
        super(standard_AE, self).__init__()  
        self.flatten = nn.Flatten()  
        self.encoder = nn.Sequential(  
            nn.Linear(784,hidden_layer),  
            nn.ReLU()  
        )  
        self.decoder =nn.Sequential(  
            nn.Linear(hidden_layer,784),  
            nn.ReLU()  
        )  
  
    def forward(self,x):  
        x = self.flatten(x)  
        encoded=self.encoder(x)  
        decoded =self.decoder(encoded)  
        return decoded, encoded
```

2.2 Loading Fashion MNIST for experiment

```
[14]: fash_test_data = datasets.FashionMNIST(root="data", train = False, download =  
      ↪True,transform = ToTensor())  
fash_test_loader = DataLoader(dataset=fash_test_data, batch_size=batch_size,  
      ↪shuffle=False)
```

2.3 Training and experiment results

```
[ ]: hid_size = [64, 128, 256]  
  
for size in hid_size:  
    model_standardAE = standard_AE(size).to(device)  
    optimizer = torch.optim.Adam(model_standardAE.parameters(), learning_rate)  
  
    train_batch_loss = []  
    train_epoch_loss = []
```

```

val_loss = []
print('Training begins for hidden unit size ', size, '!!')
print('-----')
for epoch in range(epochs):
    for batch_idx, (data, target) in enumerate(train_loader):
        data = data.to(device=device)
        target = target.to(device=device)

        model_standardAE.train()
        decoded, encoded = model_standardAE(data)
        loss = criterion(decoded, torch.flatten(data,1))
        train_batch_loss.append(loss.item())

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    train_epoch_loss.append(train_batch_loss[-1])

    model_standardAE.eval()
    with torch.no_grad():
        for (data, target) in test_loader:
            data = data.to(device)
            target = target.to(device)
            decoded, encoded = model_standardAE(data)
            val_iter_loss = criterion(decoded, torch.flatten(data,1))

    val_loss.append(val_iter_loss.item())
    print('Epoch Number: ', epoch+1)
    print('Training loss: ', train_epoch_loss[-1])
    print('Validation loss: ', val_loss[-1])
    print('-----')

fig, axs = plt.subplots(2, constrained_layout= True, figsize=(8,8))
fig.suptitle('Plots for Standard Autoencoder')
axs[0].set(ylabel='Training Loss', xlabel='Epochs')
axs[0].plot(range(1,epochs+1),train_epoch_loss, 'tab:red')
axs[1].set(ylabel='Validation Loss', xlabel='Epochs')
axs[1].plot(range(1,epochs+1),val_loss, 'tab:green')
plt.show()
print('Training done for hidden unit size ', size, '!!')
print('-----')
image = test_loader.dataset.data[1, :, :].clone() #take one image from test_
↪set
with torch.no_grad():
    image = image.view(-1,28,28).to(device=device).float()

```

```

decoded, encoded = model_standardAE.forward(image)
decoded = decoded.detach().cpu().numpy()
image = image.reshape(28,28).detach().cpu().numpy()
plt.subplot(1,2,1)
plt.imshow(image , cmap='gray')
plt.title('input image')
plt.subplots_adjust(right=1.5)
plt.subplot(1,2,2)
plt.imshow(decoded.reshape(28,28),cmap = 'gray')
plt.title("Reconstructed image using standard AE hidden unit size_
↪"+str(size))
plt.show()

fash_image = fash_test_loader.dataset.data[12, :, :].clone()
fash_image = fash_image.view(-1,28,28).to(device=device).float()
decoded, encoded = model_standardAE.forward(fash_image)
decoded = decoded.detach().cpu().numpy()
fash_image = fash_image.reshape(28,28).detach().cpu().numpy()
plt.subplot(1,2,1)
plt.imshow(fash_image , cmap='gray')
plt.title('input image')
plt.subplots_adjust(right=1.5)
plt.subplot(1,2,2)
plt.imshow(decoded.reshape(28,28),cmap = 'gray')
plt.title("Reconstructed image using standard AE hidden unit size_
↪"+str(size))
plt.show()

noisy_img = np.random.normal(loc=128,scale=10,size=(28,28))
noisy_img = torch.from_numpy(noisy_img).reshape(1,1,28,28).
↪to(device=device).float()
decoded, encoded = model_standardAE.forward(noisy_img)
decoded = decoded.detach().cpu().numpy()
noisy_img = noisy_img.reshape(28,28).detach().cpu().numpy()
plt.subplot(1,2,1)
plt.imshow(noisy_img , cmap='gray')
plt.title('input image')
plt.subplots_adjust(right=1.5)
plt.subplot(1,2,2)
plt.imshow(decoded.reshape(28,28),cmap = 'gray')
plt.title("Reconstructed image using standard AE hidden unit size_
↪"+str(size))
plt.show()
print('-----')
print('-----')

```

Training begins for hidden unit size 64 !

```
-----  
Epoch Number: 1  
Training loss: 0.013325842097401619  
Validation loss: 0.013729165308177471  
-----
```

```
Epoch Number: 2  
Training loss: 0.010477734729647636  
Validation loss: 0.012638967484235764  
-----
```

```
Epoch Number: 3  
Training loss: 0.011704877950251102  
Validation loss: 0.01246145274490118  
-----
```

```
Epoch Number: 4  
Training loss: 0.012843029573559761  
Validation loss: 0.012602227739989758  
-----
```

```
Epoch Number: 5  
Training loss: 0.011977956630289555  
Validation loss: 0.012281843461096287  
-----
```

```
Epoch Number: 6  
Training loss: 0.0108944708481431  
Validation loss: 0.011952475644648075  
-----
```

```
Epoch Number: 7  
Training loss: 0.009793129749596119  
Validation loss: 0.012324715033173561  
-----
```

3 Sparse Autoencoders

3.1 Defining sparse autoencoder class

```
[ ]: class overcomp_sparse_AE(nn.Module):  
    def __init__(self):  
        super(overcomp_sparse_AE, self).__init__()  
        self.flatten = nn.Flatten()  
        self.encoder = nn.Sequential(  
            nn.Linear(784,1521),  
            nn.ReLU())  
        self.decoder =nn.Sequential(  
            nn.Linear(1521,784),  
            nn.ReLU())
```

```

def forward(self,x):
    x = self.flatten(x)
    encoded=self.encoder(x)
    decoded =self.decoder(encoded)
    return decoded, encoded

```

3.2 Defining experimentation functions

- Average hidden layer activation avg_hl_activations
- Visualize activations visualize_activations
- Visualize learned filters encoder_decoder_filters_plots

```

[ ]: def avg_hl_activations(model,test_dataloader,model_name):
    model.eval()
    avg_act_val = 0
    with torch.no_grad():
        for (data,label) in test_dataloader:
            (data,label) = (data.to(device),label.to(device))
            decoded,encoded = model(data) #our prediction
            avg_act_val += float(torch.mean(encoded))
    avg_act_val /= len(test_dataloader)
    print("The average activation of " + str(model_name)+" is",avg_act_val)

def encoder_decoder_filters_plots(model,model_name,device):
    with torch.no_grad():
        encoder_filters = model.encoder[0].weight.detach().cpu().numpy()
        decoder_filters = model.decoder[0].weight.detach().cpu().numpy()
        #plot the encoder and decoder weights as an image for 0th neuron
        plt.imshow(encoder_filters[0].reshape(28,28), cmap='gray')
        plt.colorbar()
        plt.title('Encoder Filters for '+str(0)+'th neuron of '+
↪str(model_name))
        plt.show()

        plt.imshow(decoder_filters[:,0].reshape(28,28), cmap='gray')
        plt.colorbar()
        plt.title('Decoder Filters for '+str(0)+'th neuron of '+
↪str(model_name))
        plt.show()

def visualize_activations(model,test_dataloader,model_name,device,hidden_layer):
    ↪ #visualize the activations
    for i,ind in enumerate(idx_test):
        test_image = test_dataloader.dataset.data[ind].clone()
        test_label = test_dataloader.dataset.targets[ind].clone()
        with torch.no_grad():
            test_image = test_image.reshape(1,1,28,28).to(device=device).float()

```



```

        decoded, encoded = model.forward(test_image)
        encoded = encoded.detach().cpu().numpy()
        plt.imshow(encoded.reshape(int(np.sqrt(hidden_layer)), int(np.
→sqrt(hidden_layer))), cmap='gray')
        str_title = "Activation for digit "+str(test_label.item())
        plt.title(str_title)
        plt.show()

```

```

[ ]: sparsity_values = [0.0001, 0.001, 0.005, 0.1]

for lam in sparsity_values:
    model_over_sparseAE = overcomp_sparseAE().to(device)
    optimizer = torch.optim.Adam(model_over_sparseAE.parameters(), learning_rate)

    train_batch_loss = []
    train_epoch_loss = []
    val_loss = []
    print('Training begins for regularization lambda value ', lam, '!!')
    print('-----')
    for epoch in range(epochs):
        for batch_idx, (data, target) in enumerate(train_loader):
            data = data.to(device=device)
            target = target.to(device=device)

            model_over_sparseAE.train()
            decoded, encoded = model_over_sparseAE(data)
            loss = criterion(decoded, torch.flatten(data,1))
            loss+= lam*torch.linalg.norm(encoded, 1)
            train_batch_loss.append(loss.item())

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

        train_epoch_loss.append(train_batch_loss[-1])

    model_over_sparseAE.eval()
    with torch.no_grad():
        for (data, target) in test_loader:
            data = data.to(device)
            target = target.to(device)

            decoded, encoded = model_over_sparseAE(data)
            val_iter_loss = criterion(decoded, torch.flatten(data,1))
            val_iter_loss += lam*torch.linalg.norm(encoded, 1)

    val_loss.append(val_iter_loss.item())

```

```

print('Epoch Number: ',epoch+1)
print('Training loss: ', train_epoch_loss[-1])
print('Validation loss: ', val_loss[-1])
print('-----')

fig, axs = plt.subplots(2, constrained_layout= True, figsize=(8,8))
fig.suptitle('Plots for Standard Autoencoder')
axs[0].set(ylabel='Training Loss', xlabel='Epochs')
axs[0].plot(range(1,epochs+1),train_epoch_loss, 'tab:red')
axs[1].set(ylabel='Validation Loss', xlabel='Epochs')
axs[1].plot(range(1,epochs+1),val_loss, 'tab:green')
plt.show()
print('Training done for regularization lambda value ', lam, '!!')
print('-----')
image = test_loader.dataset.data[1, :, :].clone() #take one image from test_
↪set
with torch.no_grad():
    image = image.view(-1,28,28).to(device=device).float()
    decoded, encoded = model_over_sparseAE.forward(image)
    decoded = decoded.detach().cpu().numpy()
    image = image.reshape(28,28).detach().cpu().numpy()
    plt.subplot(1,2,1)
    plt.imshow(image , cmap='gray')
    plt.title('input image')
    plt.subplots_adjust(right=1.5)
    plt.subplot(1,2,2)
    plt.imshow(decoded.reshape(28,28),cmap = 'gray')
    plt.title("Reconstructed image using Sparse AE with lambda value "+str(lam))
    plt.show()
    avg_hl_activations(model_over_sparseAE,test_loader,"Sparse AE with lambda_
↪value "+str(lam))
    encoder_decoder_filters_plots(model_over_sparseAE,"Sparse AE with lambda_
↪value "+str(lam),device)
    visualize_activations(model_over_sparseAE,test_loader,"Sparse AE with_
↪lambda value "+str(lam),device,1521)
print('-----')
print('-----')

```

4 Denoising Autoencoders

4.1 Experimentation with different noise values

```

[ ]: noise_values = [0.3, 0.5, 0.8, 0.9]

for noise_val in noise_values:
    model_denoiseAE = standard_AE(256).to(device)

```

```

optimizer = torch.optim.Adam(model_denoiseAE.parameters(), learning_rate)

train_batch_loss = []
train_epoch_loss = []
val_loss = []
print('Training begins for noise value ', noise_val, '!!')
print('-----')
for epoch in range(epochs):
    for batch_idx, (data, target) in enumerate(train_loader):

        noise = torch.randn(data.size())*noise_val
        data = data + noise
        data = data.to(device=device)
        target = target.to(device=device)

        model_denoiseAE.train()
        decoded, encoded = model_denoiseAE(data)
        loss = criterion(decoded, torch.flatten(data,1))
        train_batch_loss.append(loss.item())

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    train_epoch_loss.append(train_batch_loss[-1])

    model_denoiseAE.eval()
    with torch.no_grad():
        for (data, target) in test_loader:
            noise = torch.randn(data.size())*noise_val
            data = data + noise
            data = data.to(device=device)
            target = target.to(device)
            decoded, encoded = model_denoiseAE(data)
            val_iter_loss = criterion(decoded, torch.flatten(data,1))

        val_loss.append(val_iter_loss.item())
    print('Epoch Number: ', epoch+1)
    print('Training loss: ', train_epoch_loss[-1])
    print('Validation loss: ', val_loss[-1])
    print('-----')

fig, axs = plt.subplots(2, constrained_layout= True, figsize=(8,8))
fig.suptitle('Plots for Standard Autoencoder')
axs[0].set(ylabel='Training Loss', xlabel='Epochs')
axs[0].plot(range(1, epochs+1), train_epoch_loss, 'tab:red')

```

```

    axs[1].set(ylabel='Validation Loss', xlabel='Epochs')
    axs[1].plot(range(1,epochs+1),val_loss, 'tab:green')
    plt.show()
    print('Training done for denoising AE with noise value',noise_val,'!')
    print('-----')
    test_image = test_loader.dataset.data[3, :, :].clone() #RANDOM TEST IMAGE
    noise = torch.randn(test_image.size())*noise_val
    image = test_image + noise
    with torch.no_grad():
        image = image.view(-1,28,28).to(device=device).float()
        decoded, encoded = model_over_sparseAE.forward(image)
        decoded = decoded.detach().cpu().numpy()
        image = image.reshape(28,28).detach().cpu().numpy()
        plt.subplot(1,2,1)
        plt.imshow(image , cmap='gray')
        plt.title('input image')
        plt.subplots_adjust(right=1.5)
        plt.subplot(1,2,2)
        plt.imshow(decoded.reshape(28,28),cmap = 'gray')
        plt.title("Reconstructed image using Denoising AE with noise value_
↪"+str(noise_val))
        plt.show()
    encoder_decoder_filters_plots( model_denoiseAE,"Denoising AE with noise_
↪value"+str(noise_val),device)
    print('-----')
    print('-----')

```